

Alexander M. Nicoara (A20301259)  
CS450: Operating Systems  
Programming Assignment 2  
Professor Leung

## Implementing *procState()*

---

In part 2 of this assignment, I was tasked with the objective of implementing a new xv6 system call *procState()*. This call prints a list of processes and some of their attributes (such as name, state, id, and memory) to the screen. It is called by the command `$ ps` in the shell (user level), and takes no arguments. In my implementation, the program (*ps.c*) makes a call to *ps()*, which generates an interrupt using the `int n` instruction, where *n* represents the system call number. Since 21 system calls already exist, *procState()* will have the number 22. After the trap is handled (we've entered the kernel), the *syscall* function (in *syscall.c*) will run and use the number *n* from the `%eax` register. The *syscall* function will call the appropriate routine (based on the value of *n*), which in our case is 22, and will call the *sys\_ps()* function in *sysproc.c*. In *sys\_ps()*, the function *ps()* will be called from *proc.c* and will execute, printing the process information for each process to the screen and then returning, which will then propagate the return in *sys\_ps()* and send us back to trap function (which called *syscall*) which will also return and then bring us back to user mode.

## Code Changes/Modifications

---

To create my system call, I first carefully analyzed other system calls (including *read*) to see how declarations were made and learn the flow/direction of how system calls get called and executed.

### In *proc.c*

- Added (starting from *line 536* until *line 561*)
  - o The reason why I am implementing my system call function here is because it has access to *ptable*, which contains a list of process attributes for each process that we need to print. This function loops through the *ptable* and prints a line displaying information (process name, state, id, memory) for each process.

```

int ps()
{
    struct proc *pr;
    sti(); // enable hardware interrupts

    acquire(&ptable.lock);
    cprintf("Name      State      ID      Memory \n"); // cprintf only prints %s, %d, %x, %p
    for(pr = ptable.proc; pr < &ptable.proc[NPROC]; pr++){
        // convert bytes to kilobytes
        int proc_mem = (int) (pr->sz);
        int memkb = proc_mem/1000; // part in front of decimal
        int mdec = proc_mem%1000; // remainder

        // print out a process state that matches the enum along with PID, Name, Memory in KB
        if(pr->state == SLEEPING){
            cprintf("%s | Sleeping | %d | %d.%d KBytes \n", pr->name, pr->pid, memkb, mdec);
        } else if(pr->state == RUNNING){
            cprintf("%s | Running | %d | %d.%d KBytes \n", pr->name, pr->pid, memkb, mdec);
        } else if(pr->state == RUNNABLE){
            cprintf("%s | Runnable | %d | %d.%d KBytes \n", pr->name, pr->pid, memkb, mdec);
        } else if(pr->state == ZOMBIE){
            cprintf("%s | Zombie | %d | %d.%d KBytes \n", pr->name, pr->pid, memkb, mdec);
        }
    }

    release(&ptable.lock);
    return 22;
}

```

Created a test program **ps.c**

- Added
  - o Before adding the system call itself, I created a test program to make a call to the *ps()* function inside of *proc.c*

```

#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"

int
main(int argc, char *argv[])
{
    ps();
    exit();
}

```

In **syscall.h**

- Added at *line 23*
  - o We define our system call and assign an integer value (22), which represents the identification number which the kernel uses to know what call to make.

```
#define SYS_ps 22
```

### In **defs.h**

- Added at *line 123*
  - o Declaration of ps()

```
int ps(void);
```

### In **user.h**

- Added at *line 26*
  - o Declaration of ps()

```
int ps(void);
```

### In **sysproc.c**

- Added at *line 93*
  - o When sys\_ps is called, execute ps() from proc.c and return 22 (the id of the sys\_ps system call). Also display a comment (that you can uncomment) to show that the function was called when debugging/testing.

```
int
sys_ps(void)
{
    // cprintf("SYSPROC.C - inside sys_ps()\n");
    return ps();
}
```

### In **usys.S**

- Added at *line 32*
  - o Assembly declaration of ps

```
SYSCALL(ps)
```

### In **syscall.c**

- Added at *line 106*
  - o Extern allows sys\_ps function to be visible from other files.

```
extern int sys_ps(void);
```

- Added (below) at *line 130*, moved closing bracket (originally line 130) down a line.
  - o This associates our system call id to be associated with the function `sys_ps` in `sysproc.c` and indexed at 22 in the IDT.

```
[SYS_ps] sys_ps,
```

- Added (below) at *line 141*, moved everything below down a line.
  - o A comment that can be uncommented to print whenever `syscall()` is called, this shows that a system call is being called.

```
// cprintf("SYSCALL.C - inside syscall()\n");
```

### In `trap.c`

- Added (below) at *line 43*, moved everything below down a line.
  - o A comment that can be uncommented to print whenever `trap()` is called, this shows that the trap handler works and is able to invoke a system call.

```
// cprintf("TRAP.C - inside trap()\n");
```

### In `Makefile`

- Added at *line 176*, shifted below lines down by 1 line.

```
_ps\
```

- Edited line 276, adding `ps.c` (the highlight represents the addition).

```
ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c ps.c zombie.c\
```

## Test Cases

Input/Expected Result	Result
<p><u>First do:</u> syscall.c – uncomment line #141 sysproc.c – uncomment line #96 trap.c – uncomment line #43 <u>Run:</u> \$ ps <u>Expectation:</u> We get print statements showing that some functions are being called up to the execution of <i>ps()</i> – This proves that my implementation is a system call.</p>	<pre>TRAP.C - inside trap() SYSCALL.C - inside syscall() SYSPROC.C - inside sys_ps() Name      State      ID      Memory init      Sleeping       1   12.288 KBytes sh      Sleeping       2   16.384 KBytes ps      Running        3   12.288 KBytes</pre>
<p><u>Run:</u> \$ ps <u>Expectation:</u> <i>ps()</i> runs and prints a list of all the processes and their attributes (name, state, id, memory)</p>	<pre>Name      State      ID      Memory init      Sleeping       1   12.288 KBytes sh      Sleeping       2   16.384 KBytes ps      Running        3   12.288 KBytes</pre>
<p><u>Open a fresh xv6 qemu window and run:</u> \$ ls then \$ echo A then \$ ps <u>Expectation:</u> The PID of <i>ps</i> is 5 (since it was 5<sup>th</sup> in sequence) and we won't see <i>ls</i> or <i>echo</i> because they terminated after execution.</p>	<pre>Name      State      ID      Memory init      Sleeping       1   12.288 KBytes sh      Sleeping       2   16.384 KBytes ps      Running        5   12.288 KBytes</pre>
<p><u>Run:</u> \$ ls   ps <u>Expectation:</u> We can see <i>ls</i> in the process table with the other processes, and <i>ls</i> is currently running. The pipe is not performing parallel execution, I used this to make the <i>ls</i> command execute before <i>ps</i>.</p>	<pre>Name      State      ID      Memory init      Sleeping       1   12.288 KBytes sh      Sleeping       2   16.384 KBytes sh      Sleeping       3   49.152 KBytes ls      Running        4   12.288 KBytes ps      Running        5   12.288 KBytes</pre>

## Final Result

```
QEMU
SeaBIOS (version Ubuntu-1.8.2-1ubuntu1)

iPXE (http://ipxe.org) 00:03.0 C980 PCI2.10 PnP PMM+1FF92460+1FED2460 C980

Booting from Hard Disk...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap s
t 58
init: starting sh
$ ps
Name      State      ID      Memory
init  | Sleeping | 1      | 12.288 KBytes
sh  | Sleeping | 2      | 16.384 KBytes
ps  | Running  | 3      | 12.288 KBytes
$
```