

Alexander M. Nicoara (A20301259)

CS450: Operating Systems

Programming Assignment 2

Professor Leung

Introduction

In part 1 of this assignment, the objective was to trace the xv6 system call *read()*. This system call has the arguments: *int fd* (the file descriptor), *void *buf* (the buffer), and *size_t n* (number of bytes to read). The *read()* system call reads from the file until *n* bytes is reached, then it returns either the number of bytes read or 0 (indicating end of the file). Upon an unsuccessful execution, -1 is returned.

Tracing *read()*

read() is first called from the user level through a program that executes the function *read(fd, 10, n)*. Next, the *usys.S* puts the *read()* system call id (22 from *syscall.h*) into the *eax* register and then generates an interrupt with *int* that is a system call.

```
in usys.S :: line 0
#include "syscall.h"
#include "traps.h"

#define SYSCALL(name)
    .globl name;
    name:
        movl $SYS_ ## name, %eax;
        int $T_SYSCALL;
```

vector.S pushes the trap number onto the stack and calls *alltraps* to context save.

```
in vector.S :: line 0
.globl vector64
vector64
    pushl $64
    jmp alltraps
```

In *trapasm.S*, the trap frame is built, data segments are pushed onto the stack, descriptors changed for memory access, and *trap(tf)* in *trap.c* is called.

```
in trapasm.S :: line 0
#include "mmu.h"

# vectors.S sends all traps here.
.globl alltraps
alltraps:
    pushl %ds
    pushl %es
```

```

pushl %fs
pushl %gs
pushal

movw $(SEG_KDATA<<3), %ax
movw %ax, %ds
movw %ax, %es

pushl %esp
call trap
addl $4, %esp

```

When we execute the function *trap()*, we check to see if the trap number references a system call. If it does (in our case), we then check to see if *myproc()* (which returns the current process) is killed. If it's not, then we set current process trapframe equal to the function's trapframe parameter. We next call the function *syscall()*, which calls *syscall()* inside of *syscall.c*

in trap.c :: line 36

```

void
trap(struct trapframe *tf)
{
    if(tf->trapno == T_SYSCALL){
        if(myproc()->killed)
            exit();
        myproc()->tf = tf;
        syscall();
    }
}

```

In the function *syscall()*, we obtain the current process from the cpu by calling *myproc()*. We obtain the value of the *eax* register and store it into *num*. If *num* is within bounds of the of the possible system call numbers, then call the function inside the *syscalls* array with index *num* and store the return value in the register *eax*.

in syscall.c :: line 131

```

void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    }
}

```

Since *num* = 22 in our case, *sys_read()* will be the function called. *sys_read()* is in *sysfile.c* and checks to see if the parameters are valid for execution (in our case they aren't since *fd* isn't valid). The first condition of the if statement tries to fetch the descriptor and the struct file.

in sysfile.c :: line 69

```

int
sys_read(void)

```

```

{
    struct file *f;
    int n;
    char *p;

    if(argfd(0, 0, &f) < 0 ||

```

argfd() is called as a result and is passed 0, 0, and the address location of f. In *argfd()*, we encounter an if statement condition which is the return value of the function *argint()*. In the condition, *argint* is called with n and &f, which returns the return value of the function *fetchint*, which ends up being 0. However, since fd doesn't contain a valid file, an exception is thrown and we return -1.

in sysfile.c :: line 21

```

static int
argfd(int n, int *pfd, struct file **pf)
{
    int fd;
    struct file *f;

    if(argint(n, &fd) < 0)

```

in syscall.c :: line 48

```

int
argint(int n, int *ip)
{
    return fetchint((myproc()->tf->esp) + 4 + 4*n, ip);

```

in syscall.c :: line 17

```

int
fetchint(uint addr, int *ip)
{
    struct proc *curproc = myproc();

    if(addr >= curproc->sz || addr+4 > curproc->sz)
        return -1;
    *ip = *(int*)(addr);
    return 0;

```

Return to argint() in sysfile.c

Return to argfd() in sysfile.c

```

    return -1;
    if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0)
        return -1;

```

Return to sys_read() in sysfile.c

Since *argfd()* returns -1 and the conditional state only needs 1 true to be true, -1 is returned and we jump back to *syscall()* and then back to *trap()*.

```

    argint(2, &n) < 0 ||

    argptr(1, &p, n) < 0)

```

```

    return -1;

..

    Return to syscall() in syscall.c
    End syscall() in syscall.c
}

..

    Return to trap() in trap.c

    if(myproc()->killed)
        exit();
    return;
}

```

Return back to trapasm.S and pop registers to restore original context. Then add trap number and error code together and returns.

```

Return to line 23 in trapasm.S
    # Return falls through to trapret...
.globl trapret
trapret:
    popal
    popl %gs
    popl %fs
    popl %es
    popl %ds
    addl $0x8, %esp    # trapno and errcode
    iret

```

We keep returning from functions and eventually get back to user mode.

```

Return to usys.S :: line 9
    ret

```